

Respite: video script

FortyGuard Hackathon'26 · live at respite.samtechpk.com · 2:32, against a hard 3:00 cap
Read the ruled lines aloud. Everything else is direction and is not spoken.

Read this first

The rules prefer a human speaking over the demo and explicitly disfavour AI narration and over-polished edits. Read it in your own voice and leave the small stumbles in. One take with a fluffed sentence beats four cuts.

The submission form caps the video at **three minutes**. This is 331 spoken words, about 2:12 of speech, plus 20 seconds watching the agent work. Read slowly it still comes in at 2:53. If a take runs close to the cap, cut the waiting around the two agent answers before cutting anything spoken.

Before you hit record

- Open <https://respite.samtechpk.com> and let the opening briefing finish writing, then reload once so it comes from cache and appears instantly. A judge should not watch a spinner.
- Check /health shows `llm_key_configured: true` and the budget is not near its limit.
- Close other tabs. Hide bookmarks. Turn off notifications.
- Have both questions copied somewhere you can paste from. Do not type them live.

What the video is for

Judging is Impact and Relevance 40%, Technical Execution 35%, Innovation 15%, Communication 10%. The finding and what a city does with it get the most time; the interface tour gets almost none. **Do not walk through features. Make one argument.**

0:00 to 0:20

The stake

On screen: the page as it loads. The headline and the three figures.

Somewhere in Phoenix tonight, someone in their seventies is lying awake in a bedroom that will not drop below twenty-eight degrees before dawn.

On the night we measured, eighteen neighbourhoods never dropped below that line. Midnight to six, not once. Fifty-eight thousand people live in them. No cool window, no recovery.

Do not say "urban heat island". Every other entry will.

0:20 to 0:52

The finding

On screen: scroll to the scatter chart. Let it sit a second. Point at the flat cloud of dots.

Cities send heat help using a social vulnerability index, assuming the most vulnerable places are also the hottest. Nobody checks it.

Here is every tract, from forty-eight thousand readings a hundred metres apart. Vulnerability along the bottom, overnight heat up the side. If that assumption held, these dots would climb left to right.

The correlation is nought point nought nought four. So fifteen neighbourhoods, fifty-two thousand people, are severely exposed and sit outside the band a vulnerability-led programme targets.

Pause. This is the moment the judges either get it or do not.

0:52 to 1:30

What a city does about it

On screen: scroll up to the console and paste the question below.

What should the city actually do tonight, and in which neighbourhoods?

Say the first line while the tool list appears, then read the quoted line off the answer. Cut the waiting in the edit: showing the tools appear is enough.

So I ask it what to do. Watch what it reads: the divergence, the tract list, the public health playbook, and its own limitations.

‘SVI alone is the wrong dispatch map.’

Most cooling centres close in the evening, before the risk period even begins. So the output is not a heat map. It is where to keep one open tonight, and where to send welfare checks.

1:30 to 2:06

What it refuses to do

On screen: paste the second question.

How many hours of relief would a cool-pavement programme buy in tract 1085.02? I need a number for a business case.

Let the refusal appear, then point at the list of sources under it.

Now the part I most want you to see. I ask for a cool-pavement effect size, for a business case.

We tested that. Control for where a tract sits, and surface composition does not explain overnight heat. So it refuses.

A model that invents a plausible figure under pressure is worse than no model, because someone will spend against it. Sixteen tests, half written to bait exactly this. All passing.

2:06 to 2:21

Our own number is soft

On screen: open the disclosure titled "Why 18 tracts, and not 32".

And our own headline is softer than it looks. Eighteen is the strictest possible reading. Allow a minute of relief to still count as none and it is thirty-two. We published that rather than hiding it.

2:21 to 2:32

Close

On screen: back to the top of the page.

Respite. Which blocks never cool down, who sleeps in them, what to do about it tonight, and where it does not know. Live at [respite dot samtechpk dot com](https://www.respite.samtechpk.com).

Reference

Every figure was read off the live site. If something on screen disagrees with this page, the site is right and this page is stale.

Pacing

Beat	Words	Speech	Watching	Ends at
The stake	51	0:20		0:20
The finding	79	0:32		0:52
What a city does about it	66	0:26	0:12	1:30
What it refuses to do	70	0:28	0:08	2:06
Our own number is soft	36	0:14		2:21
Close	29	0:12		2:32
Total	331	2:12	0:20	2:32

At a slow 130 words a minute this runs 2:53; at a brisk 170, 2:17. The cap is 3:00, so the slow figure is the one that matters. There is no stated minimum, so coming in under is not a problem.

Numbers you may be asked about

Claim	Value
Tracts with no relief	18
People in them	58,176
Tracts measured	134
Readings, at 100 m spacing	47,944
Spread, coolest to hottest block	3.70 h in one night
Afternoon spread, same area	1.57 °C
Variance between tracts (ICC)	0.855, so 86%
Exposure vs vulnerability	r = 0.004, n = 132
Severely exposed, outside the band	15 tracts, 52,091 people, 6,293 over 65
In the band, not severely exposed	53 tracts, 212,101 people
No-relief count by tolerance	18 at zero, 32 under a minute
Eval suite	16 cases, 16 clean
Land cover, controlling for position	built t = +0.24, vegetation t = -1.46

Things not to say

- "AI-powered" or "leveraging LLMs". Show the agent refusing something instead.

- **"Urban heat island."** Everyone says it, and it means nothing to a judge by the fifth entry.
- Any intervention effect size. The tool refuses to give one; do not undercut it on camera.
- **"We could add..."** Nobody scores a roadmap.
- Do not claim the cooling-centre and welfare-check actions are things we measured. They are the standard public health playbook, and the agent labels them that way on screen. Saying otherwise contradicts your own demo.
- Do not read the judging criteria back at them.

If a take goes wrong

The agent takes ten to twenty-five seconds, and the "what should the city do" answer is long. Let it run once. If you need to redo a later beat, cut back to the briefing already on screen rather than sitting through another full answer.

If the agent says you have used your allowance, you have hit the hourly cap from re-recording. It clears within the hour, or Tehseen can raise it on the box before a session.