

Respite: video script

FortyGuard Hackathon'26 · live at respite.samtechpk.com · target 4:18 of a 3 to 5 minute window

Read the ruled lines aloud. Everything else is direction and is not spoken.

Read this first

The rules prefer a human speaking over the demo and explicitly disfavour AI narration and over-polished edits. Read it in your own voice and leave the small stumbles in. One take with a fluffed sentence beats four cuts.

This is 532 spoken words, about 3:33 of speech, plus 45 seconds watching the agent work. That waiting is not dead air. It is the part that proves the thing is real.

Before you hit record

- Open <https://respite.samtechpk.com> and let the opening briefing finish writing, then reload once so it comes from cache and appears instantly. A judge should not watch a spinner.
- Check /health shows `llm_key_configured: true` and the budget is not near its limit.
- Close other tabs. Hide bookmarks. Turn off notifications.
- Have both questions copied somewhere you can paste from. Do not type them live.

What the video is for

Judging is Impact and Relevance 40%, Technical Execution 35%, Innovation 15%, Communication 10%. The finding and what a city does with it get the most time; the interface tour gets almost none. **Do not walk through features. Make one argument.**

0:00 to 0:29

The stake

On screen: the page as it loads. The headline and the three figures.

Somewhere in Phoenix tonight, someone in their seventies is lying awake in a bedroom that will not drop below twenty-eight degrees before dawn. No heatwave headline. Just a house that never cools.

On the night we measured, eighteen neighbourhoods never dropped below that line. Not once, midnight to six. Fifty-eight thousand people live in them.

No cool window, no recovery. That is why epidemiologists watch the overnight minimum, not the afternoon peak.

Do not say "urban heat island". Every other entry will.

0:29 to 0:55

What we measured

On screen: scroll to the map. Let the streets and place names show. Hover one tract.

Central Phoenix. A hundred and thirty-four neighbourhoods, from nearly forty-eight thousand readings a hundred metres apart, on the FortyGuard API.

Coollest block to hottest is three hours and forty minutes of difference in one night. By day the same city varies by a degree and a half. Uniform in the afternoon, not after dark.

And eighty-six percent of that variation is between neighbourhoods, not within them.

0:55 to 1:32

The finding

On screen: scroll to the scatter chart. Let it sit a second before speaking. Point at the flat cloud of dots.

Now the part that surprised us. Cities send heat help using a social vulnerability index, which assumes the most vulnerable places are also the hottest. Nobody checks it.

Here is every tract. Vulnerability along the bottom, overnight heat up the side. If that held, these dots would climb left to right.

The correlation is nought point nought nought four. No relationship at all, and we tried hard to break that result.

So fifteen neighbourhoods, fifty-two thousand people, six thousand over sixty-five, are severely exposed and outside the band a vulnerability-led programme targets.

Pause. This is the moment the judges either get it or do not.

1:32 to 2:32

What a city does on Monday

On screen: scroll up to the console and paste the question below.

What should the city actually do tonight, and in which neighbourhoods?

Say the first line while the tool list appears. Then stop talking, let the answer land, and read its opening sentence off the screen.

This is an agent, not a dashboard. Watch what it reads: the divergence, the tract list, the public health playbook, and its own limitations.

‘Put crews and overnight cooling access in the severely exposed tracts, not only the high-vulnerability ones. SVI alone is the wrong dispatch map.’

Most cooling centres close in the evening, before the risk period even begins. So the output is not a heat map. It is which neighbourhoods to keep one open in tonight, and where to send a finite number of welfare checks.

2:32 to 3:35

What it refuses to do

On screen: paste the second question.

How many hours of relief would a cool-pavement programme buy in tract 1085.02? I need a number for a business case.

Let the refusal render, then point at the list of sources underneath it.

Here is the part I most want you to see. Reasonable question, and we cannot answer it. We sampled fifty tracts to test whether pavement and canopy explain overnight heat. Control for where a tract sits, and they do not.

So it refuses, even when you tell it you need a number for a business case.

A model that invents a plausible figure under pressure is worse than no model, because someone will spend against it. Our limits are a tool the agent calls, not a line in a prompt, so they cannot be edited away. Sixteen tests, half written to bait exactly this. All sixteen passing.

3:35 to 4:04

Our own number is soft

On screen: open the disclosure titled "Why 18 tracts, and not 32".

One last thing, and it is what I would want to know if I were judging this. Eighteen is the strictest possible reading. Some tracts dipped below the line for a tenth of a second. Allow a minute of relief to still count as none, and it is thirty-two.

We found that in our own data, published it rather than hiding it, and the agent hands you the whole curve if you ask.

4:04 to 4:18

Close

On screen: back to the top of the page.

Respite. Which blocks never cool down, who sleeps in them, what to do about it tonight, and where it does not know.

Live now at respitedot.com. Ask it something hard.

Reference

Every figure was read off the live site. If something on screen disagrees with this page, the site is right and this page is stale.

Pacing

Beat	Words	Speech	Watching	Ends at
The stake	72	0:29		0:29
What we measured	66	0:26		0:55
The finding	92	0:37		1:32
What a city does on Monday	88	0:35	0:25	2:32
What it refuses to do	107	0:43	0:20	3:35
Our own number is soft	73	0:29		4:04
Close	34	0:14		4:18
Total	532	3:33	0:45	4:18

At a slow 130 words a minute this runs 4:51; at a brisk 170, 3:53. Both sit inside the three to five minute window, so pace is not something to worry about on the day. If you are short, slow down rather than adding material.

Numbers you may be asked about

Claim	Value
Tracts with no relief	18
People in them	58,176
Tracts measured	134
Readings, at 100 m spacing	47,944
Spread, coolest to hottest block	3.70 h in one night
Afternoon spread, same area	1.57 °C
Variance between tracts (ICC)	0.855, so 86%
Exposure vs vulnerability	$r = 0.004$, $n = 132$
Severely exposed, outside the band	15 tracts, 52,091 people, 6,293 over 65
In the band, not severely exposed	53 tracts, 212,101 people
No-relief count by tolerance	18 at zero, 32 under a minute
Eval suite	16 cases, 16 clean
Land cover, controlling for position	built $t = +0.24$, vegetation $t = -1.46$

Things not to say

- **"AI-powered"** or **"leveraging LLMs"**. Show the agent refusing something instead.
- **"Urban heat island."** Everyone says it, and it means nothing to a judge by the fifth entry.
- Any intervention effect size. The tool refuses to give one; do not undercut it on camera.
- **"We could add..."** Nobody scores a roadmap.
- Do not claim the cooling-centre and welfare-check actions are things we measured. They are the standard public health playbook, and the agent labels them that way on screen. Saying otherwise contradicts your own demo.
- Do not read the judging criteria back at them.

If a take goes wrong

The agent takes ten to twenty-five seconds, and the "what should the city do" answer is long. Let it run once. If you need to redo a later beat, cut back to the briefing already on screen rather than sitting through another full answer.

If the agent says you have used your allowance, you have hit the hourly cap from re-recording. It clears within the hour, or Tehseen can raise it on the box before a session.